

# ECE 1508: Applied Deep Learning

## Chapter 1: Preliminaries

Ali Bereyhi


`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering  
University of Toronto

Winter 2025

# Artificial Neuron

Perceptron is a *special* artificial neuron. In general, an artificial neuron is



A neuron with an  $N$ -dimensional input has  $N + 1$  learnable parameters

- $N$  weights, i.e.,  $w_1, \dots, w_N$
- a bias

From now on, when not needed, we drop them from diagram for compactness

# Artificial Neuron

The output of **neuron**  $y$  is related to its inputs as  $x$

$$y = f \left( \mathbf{w}^T x + b \right)$$

where we define

- $\mathbf{w}^T = [w_1, \dots, w_N]$  to be the vector of weights
- $b$  to be the **bias**
- $f(\cdot) : \mathbb{R} \mapsto \mathbb{R}$  to be the **activation function**

In **perceptron**, the **activation function** was *step function*<sup>1</sup>

A **neuron** can in general have any **activation**

Another special case is the linear **activation**  $f(z) = z$  by which the **neuron** reduces to the basic **linear model**


---

<sup>1</sup>It turns out soon that this was in fact a bad choice of **activation**!

# Artificial Neuron: Activation

If we intend to learn *nonlinear* functions; then, we need *nonlinear activation*

Some sample activation functions: *rectified linear unit (ReLU)* function




$$\text{ReLU}(x) = \max\{x, 0\}$$

# Artificial Neuron: Activation

If we intend to learn **nonlinear** functions; then, we need **nonlinear activation**

Some sample activation functions: **sigmoid** function




$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

# Artificial Neuron: Activation

If we intend to learn *nonlinear* functions; then, we need *nonlinear activation*

Some sample activation functions: *hyperbolic tangent* function




$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

# Artificial Neuron: Activation

If we intend to learn *nonlinear* functions; then, we need *nonlinear activation*

Some sample activation functions: *soft ReLU* function




$$\text{SReLU}(x) = \log(1 + e^x)$$

# Neural Network

## Artificial Neural Network

Artificial neural network is a **directed** graph that connects a set of **inputs** to a set of **outputs**: the nodes of this graph are **neurons** whose **activation** functions are known and whose weights and **biases** are **learnable**




Unless we are talking to biologists, we can safely drop the term **artificial** 😊



# Neural Network: *Depth*

## Depth of a Neural Network

*The longest path between an **input** and the **output***




*Here, the depth is 3*

# Deep Neural Network

## Deep Neural Network

A neural network whose *depth* is larger than 2




Here, the depth is  $3 > 2 \rightsquigarrow$  it's a deep neural network

# Neural Network: Layered Architecture


In practice, we use *neural networks* with *layered architectures*

## Layer

A subset of *neurons* that are in the same distance from *inputs*



# Neural Network: *Layered Architecture*




There are three types of layers

- **Input layer** that only contains inputs **Attention: No neuron here!**
- **Output layer** that contains the neurons computing network outputs
- **hidden layers** that contain neurons and are in between

It's readily seen from the definition that  $\text{depth} = \# \text{ hidden layers} + 1$

# Neural Network: *Layered Architecture*



It's readily seen from the definition that with **layered architecture**


$$\text{depth} = \# \text{ hidden layers} + 1$$

Hence, we could equivalently say that

*a **deep** neural network has more than one hidden layer*

# Neural Network: *Layered Architecture*

Let's try our knowledge on the **XOR** neural network




- Input layer has two inputs
- Output layer has a single neuron (perceptron)
- One hidden layer that has two neurons
  - ↳ It's **not** a **deep** neural network, since it has **only one** hidden layer
  - ↳ It's a **shallow** neural network

# Neural Networks are *Models*

- + If we talk in terms of ML components, what is a *deep neural network*?
- It's a *parameterized* function from *inputs* to *outputs*; so, it's a *model*
- + Well! What are then the *hyperparameters* and *learnable* parameters?
- Anything that describes the architecture is *hyperparameter*
  - ↳ Number of neurons, how they are connected, depth, activations, . . .
- The weights and *biases* are the *learnable* parameters
- + So, it means if the *architecture* of the neural network is known, we explicitly know which parameters we should *learn*! Right?
- Exactly! Let's look at the dummy diagram we had in previous slides!

# Neural Networks are *Models*



Here, we've chosen to have **4 neurons** with **activations**  $f_1(\cdot), \dots, f_4(\cdot)$  arranged in the **above form**: these are **hyperparameters**

Now that the **architecture** is fixed, we could say

- ↳ **Neurons 1, 2, 3** have two **inputs**: each of them has two weights and a **bias**
- ↳ **Neuron 4** has four **inputs**: it has four weights and a **bias**

So in total we have  $3 \times (2 + 1) + (4 + 1) = 14$  **learnable** parameters



# Deep Learning

- + Are we finally ready to define *Deep Learning*?
- Yes! There we go

## Deep Learning (DL)

When we use a *deep neural network* to address a learning task,  
we are doing *deep learning*

Now, let's get things a bit clear

- In *ML*, we talk about any *model*, any *loss*, any *dataset*
- In *Representation Learning*, we roughly talk about *models* that can describe nonlinear *functions*: *this includes deep neural networks*
- In *DL*, we have a *deep neural network* as the *model*

So, we can say  $DL \subset Representation\ Learning \subset ML$

# Starting with Deep Learning

*At this point, we know*

- For a given learning task, we specify the **dataset**, a **model** and a **loss**
- In DL, we use **deep neural networks** as **models**
- To accomplish the learning task, we need to train the model
  - ↳ find **weights and biases** that achieve minimal **empirical risk**

*But, we yet don't know?*

- How to minimize the **empirical loss**?
  - ↳ what **algorithm** should we use?
- What kind of **hyperparameters, i.e., architecture**, should we use?
  - ↳ how many **neurons** do we need?
  - ↳ how should we connect the **neurons**?

*This is what we learn from now on! The only last piece of preliminary knowledge that we need is the method of **gradient decent** which we study next.*