

ECE 1508S2: Applied Deep Learning

Chapter 3: Advancing Our Toolbox

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2025

Tuning Hyperparameters

A problem that we *intuitively* discussed but *left open* in Chapter 1 was

*How can we *tune* the *hyperparameters* of a model?*

In this section, we are going to *get the genie out of the bottle!*

- + Shouldn't we set *hyperparameters as much as we could?! Make NNs as deep and wide as our computer let?*
- *Not really!* We don't need *very deep and wide NNs always*
- + *But, what if we are computationally strong?! Then, we are fine! Right?!*
- *No!* In fact with *large NNs* we can do the so-called *overfitting!*

Let's see a *very simple example!*

Example: *Fitting Polynomial from Noisy Data*

We start by a *classical example* which is *not that of NN* we expect

We have a *machine* which gets *real-valued x* and *returns*

$$y = x^2 + 3x + 3$$

We however *don't know this relation*: the only thing that we know is that the *inputs* and *labels* are related *via a polynomial*

We invoke ML to *learn this machine*

Let's start with making the *ML components*, i.e.,

- 1 *Dataset*
- 2 *Model*
- 3 *Loss*

Example: Polynomial Fitting - Dataset

We start by **collecting data**: we give **input** x_b to this **machine** and measure its **output** for a **batch of inputs**. Our **measurements** are however **noisy**, i.e.,

$$v_b = x_b^2 + 3x_b + 3 + \varepsilon_b$$

where ε_b is **noise** with **bounded** magnitude, i.e., $|\varepsilon_b| \leq \alpha$ for some constant α

We make our **dataset** as

$$\mathbb{D} = \{(x_b, v_b) : i = 1, \dots, B\}$$

Example: Polynomial Fitting - Model

We know that **machine** is **polynomial**: we assume a *polynomial model*

$$y = w_0 + w_1x + w_2x^2 + \dots + w_Px^P$$

for some **integer order** P


We can write it down as

$$y = w_0 + w_1x + w_2x^2 + \dots + w_Px^P$$
$$= \underbrace{\begin{bmatrix} w_0 & w_1 & w_2 & \dots & w_P \end{bmatrix}}_{\mathbf{w}^T} \underbrace{\begin{bmatrix} x^0 \\ x^1 \\ x^2 \\ \dots \\ x^P \end{bmatrix}}_{\mathbf{h}} = \mathbf{w}^T \mathbf{h}$$

Example: Polynomial Fitting - Model

We can look at our **model** as an **NN** with **dummy neurons**

h is what we get from **hidden layer** and **w** includes **weights of output layer**



The key **hyperparameter** in this **network** is **P**

Example: Polynomial Fitting - Loss

We have a **regression** problem: *recall that*

*in **regression**, the **labels** are **real-valued***

We use **squared error** as the **loss** function, i.e.,

$$\mathcal{L}(y, v) = (y - v)^2$$

*for **measurement** v and NN's **output** y*

Now, the components are **ready**

*let's start **training***

Example: Polynomial Fitting - Training

For training, we follow what we *already learned* in previous lectures

- 1 We split \mathbb{D} into a *training dataset* and *test dataset*
- 2 We start use *gradient descent* to *train* over the *training dataset*
- 3 We *test* our *trained model* over the *test dataset*


Before we go on with *training*, let's take a look back

Naive conclusion was that *making the NN large* is *always good*; if so
when we *increase P* , we should always see *lower test risk*
However, it's *not* the case!

Let's see how the *test risk* changes against *hyperparameter P*

Example: Polynomial Fitting

Test risk against the *hyperparameter P* looks like the curve below!



Over and Underfitted Model

- + *What is happening here?*
- As we pass $P = 2$, we are **overfitting!**

Overfitting

Overfitting occurs when **training** fits the model, i.e., NN, to the **training dataset**, so that it does **not generalize** to **new data-points**

We may also **pay attention** to the term **generalize** in this definition

Generalization

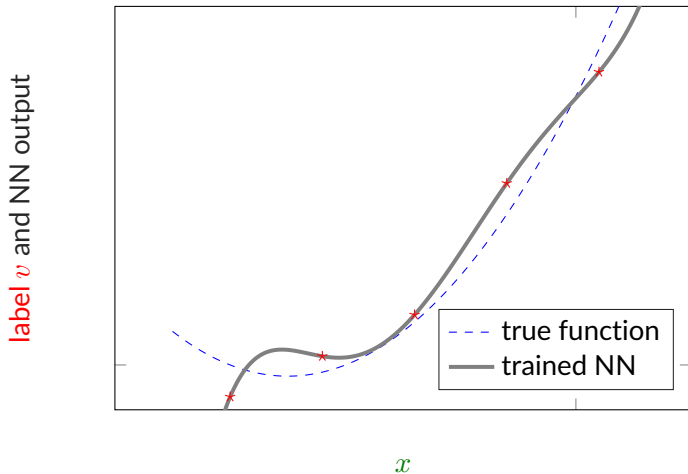
We say a **model**, i.e., **NN**, **generalizes** well if **not only its training risk**, but also **its test risk** is **small**

In simple words:

trained NN generalizes \equiv it **does what we want** on **new data**


Overfitting: Polynomial Fitting

Let's take a look back on polynomial fitting example: for large P the NN fits very well to the training data, but it deviates greatly from the true function



Overfitting: Polynomial Fitting

We can see the importance of hyperparameter tuning: if we set P to a right choice; then, our NN generalizes well, i.e., it closely track the true function



Underfitting

The **other side** of the coin is **underfitting**: it happens when **our NN does not have enough parameters** to **train**

Underfitting

Underfitting occurs when the **model**, i.e., **NN**, **neither** fits to the **training** dataset, **nor generalizes** to **new data**


We would also need to **prevent underfitting**; however,

with **current NNs**, **underfitting** can **hardly occur**

This is why it's **less discussed** in the literature

Underfitting: Polynomial Fitting

A linear model underfits in our example: setting $P = 1$ will lead to a line that can never fit our training dataset



Validation: First Step Against Overfitting

- + What is *connection* to our main task, i.e., *hyperparameter tuning*?
- Well! Before everything, we need to *tune* the *hyperparameters* right to avoid *over* or *underfitting*

Hyperparameter tuning is done by *validation*: we *change hyperparameters* among possible choices and *for each choice*, we *validate our model*


In *validation* we *train* the NN with the *specified hyperparameters* and then *test it* on a *validation set* *separated* from *training and test sets*

We set *hyperparameters* to *the choice* that

gives minimal validation risk \equiv *generalizes the best*

Validation: Polynomial Fitting

In fact *what we did* in our *dummy example* was *validation*



Why Overfitting Happens

- + But can we really do *hyperparameter tuning* in a *very deep NN*?
- **Not really!** We may *tune* some *general hyperparameters* like *number of layers*, but *cannot* really do a *complete validation*

In NNs, we invoke *other approaches* as well to combat *overfitting*

- *Regularization*
- *Dropout*
- *Data Augmentation*
- ...

To understand *these approaches*, we should first answer *the following question*

*When does **overfitting** happen in a NN?*

Let's take a look!

Why Overfitting Happens: Model Capacity

We know the **initial answer**: in our **dummy example**, it happened because we assumed **large polynomial order**

In other words

our model was **too complex** for our learning task

We can **extend this to NNs**: **overfitting** happens when **the model** is **too complex**, i.e., it's suited for learning **complicated functions**

When does **overfitting** happen in a NN? It happens when

- 1 **the model** has a **large capacity**

Though **model capacity** has a concrete definition, **for our purpose**

model capacity \equiv **ability of model** to learn **different functions**

Why Overfitting Happens: *Dataset Size*

- + But how can we *find this* out? It does *not* seem to be *easy*!
- *That's right!* This is why we look into *other reasons* as well

Let's get back to *our dummy example*: this time we check it a bit *differently*


In *our polynomial example*, we consider an *overfitted model* with $P = 5$ and *train* it on two *randomly generated datasets*

- ① a *dataset* with *20 data-points*
- ② a *dataset* with only *4 data-points*

After *training*: we *compare trained* models with the *true function*

Why Overfitting Happens: *Dataset Size*

As we can see: *overfitting* is *reduced* as we *increase the number of data-points*



Why Overfitting Happens: *Dataset Size*

This is a **general behavior**: if we have a **large enough dataset** the model **cannot really overfits** too much

- + How **large** it should be?
- It **depends on the NN**

A **general rule** is that the **more learnable parameters** the **model** has, the **larger** the **training dataset** should be

So, we can add to our answers

When does **overfitting** happen in a NN? It happens when

- ① the model has a **large capacity**
- ② our training **dataset** is **small**

Why Overfitting Happens: Co-Adaptation

Another way to see **overfitting** is to look at how *model parameters* change as *optimizer* iterates. To see it, let's get back to our *dummy polynomial-fitting NN*

Consider the following setting: we have a **high-capacity NN** with $P = 5$ and **dataset with 8 noisy samples**. We train this NN using **full-batch SGD**

We now take a look at **the trained NN** at different iterations: *recall that the vector of model parameters* is

$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_5 \end{bmatrix}$$

We start with **vector of all zeros** and keep on going

Why Overfitting Happens: Co-Adaptation

Recall that the *ground truth* \mathbf{w}^* for our polynomial machine

$$\mathbf{w}^* = \begin{bmatrix} 3 \\ 3 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Now, let's look at few iterations

$$\mathbf{w}^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_1)} = \begin{bmatrix} 2.61 \\ 2.36 \\ 0.71 \\ 0.01 \\ 0.02 \\ 0.01 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix}$$

$\mathbf{w}^{(t_2)}$ looks good! But, what if we keep on training

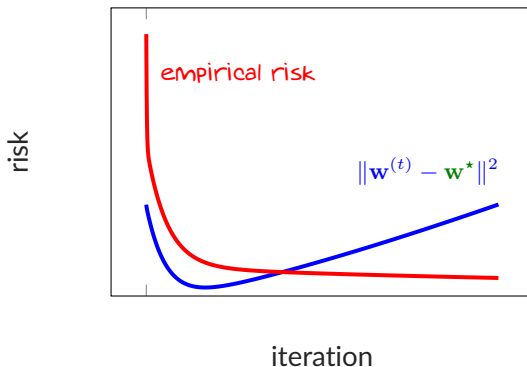
$$\mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix} \rightsquigarrow \dots \rightsquigarrow \mathbf{w}^{(t_n)} = \begin{bmatrix} 2.36 \\ 4.43 \\ 3.13 \\ -2.1 \\ 1.98 \\ -1.2 \end{bmatrix}$$

Why Overfitting Happens: Co-Adaptation

Let's formulate what we observed

Weights start to get *close to what we want* up to *intermediate number of iterations* t_2 . But, by *further training* they start to *deviate* \equiv *overfit*

We can also see this behavior in the figure below



Why Overfitting Happens: Co-Adaptation

This behavior is *co-adaptation* of the *parameters*

In initial iterations, NN fits to the *true model*: since data comes from a *quadratic function*, the *first iterations* of *SGD*

update *majorly* w_0 , w_1 and w_2


After NN has *gone* close to *ground truth*, it starts to *overfit*: due to *noise*, *quadratic model* can't *perfectly* fit; thus,

w_3 , w_4 and w_5 try to *co-adapt*, i.e., *compensate the gap* caused by *noise*

In simple words: in *first iterations* NN learns *true function*; however, at some point it starts to learn *noise*!

Co-adaptation is the *most observable implication* of *overfitting* with NNs

Co-Adaptation: Typical Learning Curve



Why Overfitting Happens: *Final List*

So, let's complete the answer list

When does overfitting happen in NNs? *It happens when*

- ① the model has a *large capacity*
- ② our training dataset is *small*
- ③ due to large number of training iterations *co-adaptation* occurs

Attention: sources are mutually related

If we have a very *large model capacity*, i.e., *very deep* with *too many neurons*; then, training it *by a small dataset* leads to *overfitting*, especially if we *keep on training for too many epochs*!

- + Now, can we do anything to avoid *overfitting*?
- Yes! Depending on what we see as *source*, we use *different tricks*

Classical Solutions to Overfitting

Overfitting happens when

- 1 the model has a **large capacity**
- 2 our training dataset is **small**
- 3 due to large number of training iterations **co-adaptation** occurs

The key tricks to address overfitting in each of these cases are

Done

- 1 We can tune the **hyperparameters** to **restrict** the **NN's capacity**
 - ↳ For instance, we can **validate** our FNN with **2, 3 and 4 hidden layers** and choose the model with **minimal validation risk**

Wait

- 2 We can **increase** our dataset by the so-called **data augmentation**
 - ↳ For instance, we can **add rotated** and **shifted** versions of **images** inside the dataset with **the same label**: a **rotated** image of a **dog** is still a **dog**!

Next

- 3 We can **regularize** our empirical risk to **penalize co-adapted solutions**
 - ↳ For instance, we can **drop out randomly** some **neurons** in each **mini-batch**

Training by Penalized Risk: *Regularization*

Regularization aims to resolve *overfitting* by treating *co-adaptation*

Let's recall *co-adaptation* in our *dummy polynomial fitting NN*: we set $P = 5$ and train our NN via the *noisy samples* inside *training dataset*; clearly,

as *training* proceeds, *empirical risk drops*

In our particular example with

$$\mathbf{w}^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_2)} = \begin{bmatrix} 3.03 \\ 2.97 \\ 0.98 \\ 0.21 \\ 0.12 \\ 0.09 \end{bmatrix} \rightsquigarrow \mathbf{w}^{(t_n)} = \begin{bmatrix} 2.36 \\ 4.43 \\ 3.13 \\ -2.1 \\ 1.98 \\ -1.2 \end{bmatrix}$$

This means that $\hat{R}(\mathbf{w}^{(t_n)}) \leq \hat{R}(\mathbf{w}^{(t_2)}) \leq \hat{R}(\mathbf{w}^{(0)})$

Training by Penalized Risk: *Regularization*

Regularization follows **this idea**: can we **modify empirical** risk, such that it **stops dropping** after t_2 ?

Let's continue with **our example**: assume **risk's value** at each \mathbf{w} is

$$\hat{R}(\mathbf{w}^{(t_3)}) = 0.001 \quad \hat{R}(\mathbf{w}^{(t_2)}) = 0.01 \quad \hat{R}(\mathbf{w}^{(0)}) = 100$$

We may note that as **training progresses**, vector $\mathbf{w}^{(t)}$ becomes **larger**. So, what if we add a **penalty** to **risk** that is **proportional to $\|\mathbf{w}^{(t)}\|^2$** : this way when **risk becomes too small**, this **penalty becomes large** and thus the **sum increases**. Let's look at this **sum** at *different iterations*

$$\tilde{R}(\mathbf{w}^{(0)}) = \hat{R}(\mathbf{w}^{(0)}) + \|\mathbf{w}^{(0)}\|^2 = 100$$

$$\tilde{R}(\mathbf{w}^{(t_2)}) = \hat{R}(\mathbf{w}^{(t_2)}) + \|\mathbf{w}^{(t_2)}\|^2 = 19.04$$

$$\tilde{R}(\mathbf{w}^{(t_n)}) = \hat{R}(\mathbf{w}^{(t_n)}) + \|\mathbf{w}^{(t_n)}\|^2 = 44.761$$

Training by Penalized Risk: Regularization

Penalized risk shows a different behavior

$$\tilde{R}(\mathbf{w}^{(0)}) = 100 \quad \tilde{R}(\mathbf{w}^{(t_2)}) = 19.04 \quad \tilde{R}(\mathbf{w}^{(t_n)}) = 44.761$$

From above values, we can say: if we apply **SGD** to minimize **penalized risk** we **may get** from $\mathbf{w}^{(0)}$ to $\mathbf{w}^{(t_2)}$; however, we will **not** get from $\mathbf{w}^{(t_2)}$ to $\mathbf{w}^{(t_n)}$

This idea is called **regularization** which can **prevent** NNs from **overfitting**

Regularization

In **training** with **regularization**, we minimize a **penalized (regularized)** form of the empirical risk, i.e.,

$$\min_{\mathbf{w}} \hat{R}(\mathbf{w}) + \Pi(\mathbf{w}) \quad (\text{Regularized Training})$$

$\Pi(\mathbf{w})$ is a **penalty** that describes the behavior of \mathbf{w} in the case of **overfitting**

Classical Regularization Approaches

There are various regularization penalties: some important ones are

- ℓ_2 or Tikhonov regularization in which we add a term proportional to $\|\mathbf{w}\|^2$

$$\Pi(\mathbf{w}) = \lambda \|\mathbf{w}\|^2$$

↳ This way we avoid very large weights

↳ This prevents perfect fit to training dataset reducing chance of overfitting

- ℓ_1 or Lasso regularization in which we add a term proportional to $\|\mathbf{w}\|_1$

$$\Pi(\mathbf{w}) = \lambda \|\mathbf{w}\|_1 = \lambda \sum_{i=1}^D |w_i|$$

↳ This way enforce \mathbf{w} to be sparse, i.e., to have to many zeros

↳ This way we reduce the capacity of NN and thus prevent overfitting

Regularizing by *Dropout*


A less conventional **regularization** approach is **dropout** that was proposed by Hinton et al. first in *their 2012 paper* and then in *their 2014 paper*:¹ the idea is at the same time **easy** and **effective**

for each **training** iteration, we **deactivate** some **nodes** of NN **at random**
or in other words we **drop them out**

¹Click to check out the papers!


Dropout: Schematic

Let's say this is the *dense* NN




Dropout: Schematic

For *first forward-backward* we select few *nodes* in *each layer*




Dropout: Schematic

For *second forward-backward* we select new *nodes* in *each layer* at random




Dropout: Schematic

For *next forward-backward* we select again few *nodes* in *each layer* at random



Dropout: Schematic

At the end, we *average* the gradients determined over these *reduced NNs*



Dropout: *Intuition*

- + But, why does *dropout* work?
- We can explain it *heuristically*

Recall that observing *overfitting* means that *NN* is *larger* than *required*. With *dropout*, in each iteration we train a *smaller version* of *NN*

- We *randomly* switch among these *smaller versions*
- Many of these *smaller versions* *do not* *overfit*

We can look at the training loop with *dropout* as an averaged training of these *smaller NNs*; hence, training loop gets *less chances* to *overfit*

- + Why do we do it *randomly*? Why not *sticking* to one *smaller version*?
- *Not* all *smaller NNs* are *good*, and we *cannot* check all of them: it's *exponentially hard* to *check all* *smaller NNs*

Dropout: Training Loop

How does **training** change with **dropout**? Training with **dropout** is exactly as before. Say we use **mini-batches**: for each **mini-batch**

- we compute the **gradient** by **forward** and **backpropagation**
- we give the **gradient** to the **optimizer** to apply the next **iteration**

The **only thing** that is **different** now is that we set the **output** of some nodes to **zero** in the **forward** pass \equiv only **forward** propagation **changes**

Let's make it concrete: when we pass **forward**, we generate **random masks** for each **layer** $\ell = 0, \dots, L$. Mask of **layer** ℓ is a vector whose length is **layer's width** and **entries are 0 or 1**, i.e., $\mathbf{s}_\ell \in \{0, 1\}^{\mathcal{W}_\ell}$. Entries of \mathbf{s}_ℓ are generated **randomly**

$$\text{each entry of } \mathbf{s}_\ell = \begin{cases} 1 & \text{with probability } p_\ell \\ 0 & \text{with probability } 1 - p_\ell \equiv \text{dropout probability} \end{cases}$$

Dropout: *Forward Propagation*

Let's show generation of **random mask** \mathbf{s}_ℓ by following notation

$$\mathbf{s}_\ell = \text{mask}(\mathcal{W}_\ell | p_\ell)$$

We are going to do **forward** propagation for each **data-point** as

DropoutForwardProp():

```

1: Initiate with  $\mathbf{y}_0 = \mathbf{x}$ 
2: for  $\ell = 0, \dots, L$  do
3:   Generate  $\mathbf{s}_\ell = \text{mask}(\mathcal{W}_\ell | p_\ell)$  # random mask
4:   Set  $\mathbf{y}_\ell = \mathbf{y}_\ell \odot \mathbf{s}_\ell$  # dropout nodes
5:   Add  $\mathbf{y}_\ell[0] = 1$  and determine  $\mathbf{z}_{\ell+1} = \mathbf{W}_{\ell+1}\mathbf{y}_\ell$  # forward affine
6:   Determine  $\mathbf{y}_{\ell+1} = f_{\ell+1}(\mathbf{z}_{\ell+1})$  # forward activation
7: end for
8: for  $\ell = 1, \dots, L + 1$  do
9:   Return  $\mathbf{y}_\ell$  and  $\mathbf{z}_\ell$ 
10: end for
```

Dropout: Backpropagation

The backpropagation goes **exactly as before**: of course those **outputs** that were **dropped out** participate with value **zero** in gradient computation

One final **piece of trick**

After **training** is over, we **scale weights** of **each layer** with its **retain probability** p_ℓ : say T is the **last iteration** of training loop; then, we finally do

$$\mathbf{W}_\ell^{(T)} \leftarrow p_\ell \mathbf{W}_\ell^{(T)}$$

- + Why do we do that?
- Well! It's **practically** understood; however, we can justify it as follows: each weight could be **what has been computed** with probability p_ℓ and **zero** with probability $1 - p_\ell$. We hence compute the **average**

Dropout: Implementation

Dropout is implemented in almost all *deep learning libraries*

```
>> import torch  
>> torch.nn.Dropout()
```

Typical choices of *retain probability* p_ℓ are

- for *input layer*, i.e., *layer 0*, $p_\ell = 0.8$
- for *hidden layers* $p_\ell = 0.5$

It's generally suggested to *drop out* more at *hidden layers*

With dropout forward pass changes in *training* and *evaluation*

- In *training* we use random mask \rightsquigarrow `model.train()`
- In *evaluation* we *don't* use random mask \rightsquigarrow `model.eval()`