

ECE 1508: Applied Deep Learning

Chapter 4: Convolutional Neural Networks

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2025

Big Picture: What to Train

In CNNs we need to train all *learnable parameters*, i.e.,

- *weights* and *biases* of output FNN
- *weights* in the *filters* of convolutional layers
- if we use *advanced pooling* with *weights*; then, we should find them as well

To train, we get dataset $\mathbb{D} = \{(\mathbf{X}_b, \mathbf{v}_b) : b = 1, \dots, B\}$ and train the CNN as


$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\boldsymbol{\theta}_b, \mathbf{v}_b) \quad (\text{Training})$$

where $\boldsymbol{\theta}_b = \text{CNN}(\mathbf{X}_b | \mathbf{w})$ for the tensor-type data-point \mathbf{X}_b with *label* \mathbf{v}_b

- ↳ We solve this optimization via a *gradient-based method*
- ↳ This means we should be able to pass *forward* and *backward*

Computation Graph

Computation graph for single data-point \mathbf{X} and its true label \mathbf{v} is as follows



For *simplicity*, let's assume that we are doing *basic SGD*

- ↳ We need to pass *forward* \mathbf{X} over this graph to get *loss*
- ↳ We should then pass *backward* to compute *gradient*

Let's try both directions

Forward Pass over CNN


The forward pass is what we learned in the last section

- We pass \mathbf{X} through first convolution to get \mathbf{Z}_1
- We **activate** \mathbf{Z}_1 to get \mathbf{Y}_1
- We pool entries of \mathbf{Y}_1 and get $\hat{\mathbf{Y}}_1$
- \vdots
- We flatten and pass **forward** through the **output FNN**

Once we are over, we have

- ↳ all **convolution**, **activated** and **pooled** values in convolutional layers
- ↳ all **affine** and **activated** values in the FNN

Backward Pass



We now want to **backpropagate**

↳ We first compute $\nabla_{\theta} \hat{R}$

↳ We **backpropagate** over the FNN **till we get to its input**


At this point, we sort $\nabla_{\hat{y}_L} \hat{R}$ in tensor by **reversing the flattening**

we now have $\nabla_{\hat{Y}_L} \hat{R}$

We only need to learn how to **backpropagate** through **pooling** and **convolutional** layers and then we can **complete** the backward pass!

Backpropagate through Pooling: Max-Pooling

Let's start with a **simple example**: in the following **pooling layer**, we have partial derivatives with respect to **pooled variables** and want to compute the partial derivatives with respect to **input of pooling layer**



Backpropagation: Max-Pooling

We can use **chain rule**

$$\frac{\partial \hat{R}}{\partial y_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_1} = 0$$

$$\frac{\partial \hat{R}}{\partial y_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_2} = \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{\partial \hat{R}}{\partial \hat{y}_2}$$

$$\frac{\partial \hat{R}}{\partial y_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_3} = 0$$

⋮

Let's see its visualization


$$\frac{\partial \hat{R}}{\partial y_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_7} = \frac{\partial \hat{R}}{\partial \hat{y}_3}$$

$$\frac{\partial \hat{R}}{\partial y_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_8} = 0$$

$$\frac{\partial \hat{R}}{\partial y_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_9} = \frac{\partial \hat{R}}{\partial \hat{y}_4}$$


Backpropagation: Max-Pooling

Each derivative in **green map** gets multiplied with its filter and added to corresponding entries on **blue map**



Backpropagate through Pooling: *Mean-Pooling*

Lets now consider **mean-pooling**: we have partial derivatives with respect to **pooled variables** and want to compute the partial derivatives with respect to **input of pooling layer**



Backward Pass: Mean-Pooling

We again use [chain rule](#)

$$\frac{\partial \hat{R}}{\partial y_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_1} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1}$$

$$\frac{\partial \hat{R}}{\partial y_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_2} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_2}$$

$$\vdots$$


$$\frac{\partial \hat{R}}{\partial y_5} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_5} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_1} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_2} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_3} + \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_4}$$

$$\vdots$$

$$\frac{\partial \hat{R}}{\partial y_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial y_9} = \frac{1}{4} \frac{\partial \hat{R}}{\partial \hat{y}_4}$$

Backward Pass: Mean-Pooling


It has same visualization: the **filter** is however this time fixed



But we can visualize it **even better!**

Backward Pass: Mean-Pooling

In fact, this is simply a *convolution*



Backpropagation through Pooling: *Summary*

Depending on **pooling** function, backpropagation can be different

↳ *It's however usually described by a **convolution***

Moral of Story

Backpropagation through **pooling** can be done by a **convolution-type operation**

Backpropagation through Convolution: Activation

Convolutional layer has two operations

↳ linear *convolution*

↳ entry-wise *activation*

The entry-wise *activation* is readily backpropagate: say we have

$$\mathbf{Y} = f(\mathbf{Z})$$


Then, we can *backpropagate* as in the *fully-connected* FNNs

$$\nabla_{\mathbf{Z}} \hat{R} = \nabla_{\mathbf{Y}} \hat{R} \odot \dot{f}(\mathbf{Z})$$

Now, let's look at linear convolution!

Backpropagation: 2D Convolution

Let's find it out through a **simple example**: in the following **convolutional layer**, we have partial derivatives with respect to **convolved variables** and want to compute the partial derivatives with respect to **input of convolutional layer**



Backpropagation: 2D Convolution

Let's start with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = w_1 \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = w_2 \frac{\partial \hat{R}}{\partial z_1} + w_1 \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = w_2 \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = w_3 \frac{\partial \hat{R}}{\partial z_1} + w_1 \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_5} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_5} = w_4 \frac{\partial \hat{R}}{\partial z_1} + w_3 \frac{\partial \hat{R}}{\partial z_2} + w_2 \frac{\partial \hat{R}}{\partial z_3} + w_1 \frac{\partial \hat{R}}{\partial z_4}$$

Backpropagation: 2D Convolution

We can keep on till finish with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_6} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_6} = w_4 \frac{\partial \hat{R}}{\partial z_2} + w_2 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial x_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_7} = w_3 \frac{\partial \hat{R}}{\partial z_3}$$


$$\frac{\partial \hat{R}}{\partial x_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_8} = w_4 \frac{\partial \hat{R}}{\partial z_3} + w_3 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial x_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_9} = w_4 \frac{\partial \hat{R}}{\partial z_4}$$

Let's look at the visualization

Backpropagation: 2D Convolution


We can use the same visualization as in pooling



Or even a *better* visualization!

Backpropagation: 2D Convolution

It's again simply a *convolution*



with *filter* being *reversed up-to-down* and *right-to-left*

Backpropagation: 2D Convolution

To backpropagate a convolutional layer, we convolve the output gradient with the **filter** being **reversed up-to-down** and **right-to-left**. To match the dimension, we simply apply **zero-padding**

Backpropagation through 2D Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W})$, where \mathbf{X} is the **input map**, i.e., a matrix, \mathbf{W} is **filter** and \mathbf{Z} is the **output map**. Assume that we know the gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to the **input map** is

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}} \right)$$


where $\check{\mathbf{W}}$ is the **up-to-down** and **left-to-right** reverse of **filter** \mathbf{W}

Backpropagation: 2D Convolution

- + What if the dimensions do **not** match?
- The dimensions can **only not match** if
 - 1 we do the forward convolution with a **stride different from one**
 - ↳ We said that this is simply **resampling**
 - ↳ We are going to **discuss it later**
 - ↳ For now assume that **we do the convolution with stride one**
 - 2 we did **no zero-padding** in the forward convolution
 - ↳ We simply **do zero-padding** in backward convolution to match the dimensions
- + What about the **gradient** with respect to **filter itself**? Don't we need it?
- Let's check it out

Backpropagation: Convolution Filter

We try it for **our example**: we have partial derivatives with respect to **convolved variables** and want to compute the partial derivatives with respect to **weights in the filter**



Backpropagation: Convolution Filter

As always, we use *chain rule*

$$\frac{\partial \hat{R}}{\partial w_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_1} = x_1 \frac{\partial \hat{R}}{\partial z_1} + x_2 \frac{\partial \hat{R}}{\partial z_2} + x_4 \frac{\partial \hat{R}}{\partial z_3} + x_5 \frac{\partial \hat{R}}{\partial z_4}$$


$$\frac{\partial \hat{R}}{\partial w_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_2} = x_2 \frac{\partial \hat{R}}{\partial z_1} + x_3 \frac{\partial \hat{R}}{\partial z_2} + x_5 \frac{\partial \hat{R}}{\partial z_3} + x_6 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial w_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_3} = x_4 \frac{\partial \hat{R}}{\partial z_1} + x_5 \frac{\partial \hat{R}}{\partial z_2} + x_7 \frac{\partial \hat{R}}{\partial z_3} + x_8 \frac{\partial \hat{R}}{\partial z_4}$$

$$\frac{\partial \hat{R}}{\partial w_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial w_4} = x_5 \frac{\partial \hat{R}}{\partial z_1} + x_6 \frac{\partial \hat{R}}{\partial z_2} + x_8 \frac{\partial \hat{R}}{\partial z_3} + x_9 \frac{\partial \hat{R}}{\partial z_4}$$

Backpropagation: Convolution Filter

It's another *convolution* with $\nabla_{\mathbf{z}} \hat{R}$ being the filter!



Backpropagation: Convolution Filter

Once we backpropagate to the **output** of a **convolutional layer**, then we can find the gradient with respect to convolution filter by **convolving output gradient** with the **input map**

Gradient w.r.t. Filters

Let $\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W})$, where \mathbf{X} is the **input map**, i.e., a matrix, \mathbf{W} is **filter** and \mathbf{Z} is the **output map**. Assume that we know gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, gradient of loss \hat{R} with respect to the **filter** is

$$\nabla_{\mathbf{W}} \hat{R} = \text{Conv}(\mathbf{X} | \nabla_{\mathbf{Z}} \hat{R})$$

- + But we checked only **2D case**! What about **multi-channel case**?!
 - Well, we can simply do it by **chain rule**

Backpropagation: Multi-Channel Convolution


Lets start with a simple case: we have a C -channel input, i.e., a tensor input, \mathbf{X} and we compute a single feature map, i.e., a matrix, \mathbf{Z}

↳ Filter \mathbf{W} has also C channels

Let's denote channel c of input and filter by $\mathbf{X}[c]$ and $\mathbf{W}[c]$; then, we can write


$$\mathbf{Z} = \sum_{c=1}^C \text{Conv}(\mathbf{X}[c] | \mathbf{W}[c])$$

Backpropagation: *Multi-Channel Convolution*



Backpropagation: Multi-Channel Convolution

Channel c of **input** is connected to the **output** map by a 2D convolution



So, we could say

$$\nabla_{\mathbf{X}[c]} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[c] \right)$$

Backpropagation: Multi-Channel Convolution

To backpropagate a convolutional layer with **tensor input** and **matrix** output, we convolve the output gradient with the **filter of each channel** being **reversed up-to-down** and **right-to-left**


Backpropagation through 2D Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W})$, where \mathbf{X} is a C -channel **input tensor**, \mathbf{W} is C -channel **filter** and \mathbf{Z} is a single-channel **output map**, i.e., **a matrix**. Assume that we know the gradient of loss \hat{R} with respect to the **output map**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to **input tensor** is

$$\nabla_{\mathbf{X}} \hat{R} = \left[\text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[1] \right), \dots, \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \check{\mathbf{W}}[C] \right) \right]$$

Note that $\nabla_{\mathbf{X}} \hat{R}$ is also a C -channel tensor

Backpropagation: *Multi-Channel Convolution*



Backpropagation: Multi-channel Convolution

We now go for the general case: we have a C -channel input tensor \mathbf{X} and we compute K -channel feature tensor \mathbf{Z}


↳ We have K filters $\mathbf{W}_1, \dots, \mathbf{W}_K$

↳ Each filter has C channels

Let's denote channel k of the output by \mathbf{Z}_k ; then, we can write

$$\mathbf{Z}_k = \text{Conv}(\mathbf{X} | \mathbf{W}_k)$$

Backpropagation: *Multi-Channel Convolution*



Backpropagation: Multi-channel Convolution

We write the **chain rule** with a bit **cheating**: let's denote the **derivative** of an object A with respect to object B as $\nabla_B A$

- ↳ if A and B are both scalars then its simple **derivative**
- ↳ if A is a scalar and B is a vector it's **gradient**
- ↳ if A and B are both vectors then its **Jacobian**
- ↳ ...

When $\mathbf{Z} = f(\mathbf{X})$ and $\hat{R} = g(\mathbf{Z})$: **chain rule** says


$$\nabla_{\mathbf{X}} \hat{R} = \nabla_{\mathbf{Z}} \hat{R} \circ \nabla_{\mathbf{X}} \mathbf{Z}$$

where \circ is a kind of product

We now write the **chain rule** with this simplified notation

Backpropagation: Multi-Channel Convolution

We do know $\nabla_{\mathbf{Z}} \hat{R} = [\nabla_{\mathbf{Z}[1]} \hat{R}, \dots, \nabla_{\mathbf{Z}[K]} \hat{R}]$



Backpropagation: Multi-Channel Convolution

The output tensor can be seen as K functions of the **input tensor**

$$\mathbf{Z}[1] = f_1(\mathbf{X}) \quad \dots \quad \mathbf{Z}[K] = f_K(\mathbf{X})$$

So, the **chain rule** can be written as

$$\begin{aligned} \nabla_{\mathbf{X}} \hat{R} &= \sum_{k=1}^K \underbrace{\nabla_{\mathbf{Z}[k]} \hat{R} \circ \nabla_{\mathbf{X}} \mathbf{Z}[k]}_{\text{what we calculated for single output map}} \\ &= \sum_{k=1}^K \left[\text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right), \dots, \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right] \\ &= \left[\sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right), \dots, \sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right] \end{aligned}$$


Backpropagation: Multi-Channel Convolution

The backward pass is therefore

$$\nabla_{\mathbf{X}} \hat{R} = \left[\sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[1] \right) \dots \sum_{k=1}^K \text{Conv} \left(\nabla_{\mathbf{Z}[k]} \hat{R} | \check{\mathbf{W}}_k[C] \right) \right]$$

Let's look at a particular **input channel c**

Backpropagation: Multi-Channel Convolution



Backpropagation: Multi-Channel Convolution

This is a new multi-channel convolution: let's define *K-channel filter* \mathbf{W}_c^\dagger for $c = 1, \dots, C$ as follows

$$\mathbf{W}_c^\dagger = \left[\check{\mathbf{W}}_1[c] \dots \check{\mathbf{W}}_K[c] \right]$$

Then, channel c of $\nabla_{\mathbf{X}} \hat{R}$ is the convolution of $\nabla_{\mathbf{Z}} \hat{R}$ with \mathbf{W}_c^\dagger


$$\nabla_{\mathbf{X}[c]} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_c^\dagger \right)$$

Or shortly, we can write

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_1^\dagger, \dots, \mathbf{W}_C^\dagger \right)$$

Backpropagation: Multi-Channel Convolution

This means that we can look at channel c of $\nabla_{\mathbf{x}} \hat{R}$ as



Backpropagation: Multi-Channel Convolution

To backpropagate to channel c of **tensor input**, we convolve the output gradient **tensor** with the K -channel **filter tensor** that is constructed by collecting the channel c of all K forward filters, each being **reversed up-to-down** and **right-to-left**

Backpropagation through Multi-Channel Convolution

Let $\mathbf{Z} = \text{Conv}(\mathbf{X} | \mathbf{W}_1, \dots, \mathbf{W}_K)$, where \mathbf{X} is a C -channel **input tensor**, \mathbf{W}_k is C -channel **filter** and \mathbf{Z} is a K -channel **output tensor**. Assume that we know the gradient of loss \hat{R} with respect to the **output tensor**, i.e., $\nabla_{\mathbf{Z}} \hat{R}$. Then, the gradient of loss \hat{R} with respect to **input tensor** is

$$\nabla_{\mathbf{X}} \hat{R} = \text{Conv} \left(\nabla_{\mathbf{Z}} \hat{R} | \mathbf{W}_1^\dagger, \dots, \mathbf{W}_K^\dagger \right)$$

where $\mathbf{W}_c^\dagger = \left[\check{\mathbf{W}}_1[c] \dots \check{\mathbf{W}}_K[c] \right]$ is a K -channel filter

Backpropagation through Convolution: Summary

Moral of Story

We can always backpropagate through a **convolutional layer** with C input channels and K output channels by

- 1 multiply output gradient entry-wise with derivative of **activation function**
- 2 convolve **output gradient** with C different **K -channel filters** computed by **rearranging of the K forward filters**

To compute the gradient with respect to **weights of channel c in filter k**

↳ **convolve** channel c of input tensor with channel k of output gradient

- + Sounds great! But what about cases with **stride $\neq 1$!**
- Well! we said we can decompose them as convolution/pooling with **stride 1** + a **resampling unit**. We just need to learn how to backpropagate through a **resampling unit**

Backpropagation: *Downsampling*

Let's again try an **example**: we have partial derivatives with respect to **down-sampled variables** and want to compute the partial derivatives with respect to **input variables**

x_1	x_2	x_3
x_4	x_5	x_6
x_7	x_8	x_9

$$\mathbf{Z} = \text{dSample}(\mathbf{X}|2)$$

~~~~~>

|             |             |
|-------------|-------------|
| $z_1 = x_1$ | $z_2 = x_3$ |
| $z_3 = x_7$ | $z_4 = x_9$ |

# Backpropagation: *Downsampling*

Let's write with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = 0$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = \frac{\partial \hat{R}}{\partial z_2}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = 0$$

⋮

$$\frac{\partial \hat{R}}{\partial x_7} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_7} = \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_8} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_8} = 0$$

$$\frac{\partial \hat{R}}{\partial x_9} = \sum_{i=1}^4 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_9} = \frac{\partial \hat{R}}{\partial z_4}$$

# Backpropagation: *Downsampling*

This is simply an *upsampling* with the same factor

|                                         |   |                                         |
|-----------------------------------------|---|-----------------------------------------|
| $\frac{\partial \hat{R}}{\partial z_1}$ | 0 | $\frac{\partial \hat{R}}{\partial z_2}$ |
| 0                                       | 0 | 0                                       |
| $\frac{\partial \hat{R}}{\partial z_3}$ | 0 | $\frac{\partial \hat{R}}{\partial z_4}$ |

uSample ( $\nabla_{\mathbf{z}} \hat{R} | 2$ )

|                                         |                                         |
|-----------------------------------------|-----------------------------------------|
| $\frac{\partial \hat{R}}{\partial z_1}$ | $\frac{\partial \hat{R}}{\partial z_2}$ |
| $\frac{\partial \hat{R}}{\partial z_3}$ | $\frac{\partial \hat{R}}{\partial z_4}$ |

## Backpropagation: Upsampling

Now, let's look into **upsampling**: we have *partial derivatives with respect to up-sampled variables* and want to compute the *partial derivatives with respect to input variables*

|             |           |             |
|-------------|-----------|-------------|
| $z_1 = x_1$ | $z_2 = 0$ | $z_3 = x_2$ |
| $z_4 = 0$   | $z_5 = 0$ | $z_6 = 0$   |
| $z_7 = x_3$ | $z_8 = 0$ | $z_9 = x_4$ |

$$\mathbf{Z} = \text{uSample}(\mathbf{X}|2)$$

←

|       |       |
|-------|-------|
| $x_1$ | $x_2$ |
| $x_3$ | $x_4$ |

## Backpropagation: *Downsampling*

Let's write with **chain rule**

$$\frac{\partial \hat{R}}{\partial x_1} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_1} = \frac{\partial \hat{R}}{\partial z_1}$$

$$\frac{\partial \hat{R}}{\partial x_2} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_2} = \frac{\partial \hat{R}}{\partial z_3}$$

$$\frac{\partial \hat{R}}{\partial x_3} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_3} = \frac{\partial \hat{R}}{\partial z_7}$$

$$\frac{\partial \hat{R}}{\partial x_4} = \sum_{i=1}^9 \frac{\partial \hat{R}}{\partial z_i} \frac{\partial z_i}{\partial x_4} = \frac{\partial \hat{R}}{\partial z_9}$$

# Backpropagation: Upsampling

This is **downsampling** with the same factor

|                                         |                                         |                                         |
|-----------------------------------------|-----------------------------------------|-----------------------------------------|
| $\frac{\partial \hat{R}}{\partial z_1}$ | $\frac{\partial \hat{R}}{\partial z_2}$ | $\frac{\partial \hat{R}}{\partial z_3}$ |
| $\frac{\partial \hat{R}}{\partial z_4}$ | $\frac{\partial \hat{R}}{\partial z_5}$ | $\frac{\partial \hat{R}}{\partial z_6}$ |
| $\frac{\partial \hat{R}}{\partial z_7}$ | $\frac{\partial \hat{R}}{\partial z_8}$ | $\frac{\partial \hat{R}}{\partial z_9}$ |

dSample  $\left( \nabla_{\mathbf{z}} \hat{R} | 2 \right)$

~~~~~→

| | |
|---|---|
| $\frac{\partial \hat{R}}{\partial z_1}$ | $\frac{\partial \hat{R}}{\partial z_3}$ |
| $\frac{\partial \hat{R}}{\partial z_7}$ | $\frac{\partial \hat{R}}{\partial z_9}$ |

Backpropagation: *Resampling*

Backpropagation through Downsampling

To backpropagate through a **downsampling** unit with factor (stride) S we **up-sample** the output gradient with factor S

Backpropagation through Upsampling

To backpropagate through an **upsampling** unit with factor (stride) S we **down-sample** the output gradient with factor S

Forward and Backpropagation in CNNs: Summary

To each **forward** action, there is a **backward** counterpart

In **forward** pass we do

- **forward** FNN
- **pooling**
- **convolution**
- **upsampling**
- **downsampling**

In **backward** pass we do

- **backward** FNN
- **backward pooling** ~ **convolution**
- **convolution** with **reversed filters**
- **downsampling**
- **upsampling**

Once we over with backward pass, we compute **all required gradients** from **forward** and **backward** variables

We Use Advanced Techniques

For FNN, we looked into advanced techniques such as

- *Dropout* and *Regularization*
 - ↳ to reduce the impact of *overfitting*
- *Input* Normalization
 - ↳ to improve training with *unbalanced* datasets
- *Batch* Normalization
 - ↳ to make the training more stable against *feature variations*
- *Data Preprocessing*
 - ↳ to *clean* our training dataset

Same goes with CNNs

we should use all these methods for same purposes in CNNs

Other Forms of Convolution

The convolution operation we considered in this chapter is often called


2D Convolution

because it screens *input* in a *2D* fashion, i.e., *left-to-right* and *up-to-down*

2D convolution is the *most popular form* of convolution; however,



- we could have *1D convolutions*
 - ↳ The filter *only slides in one direction*, i.e., *left-to-right* or *up-to-down*
 - ↳ This is useful with *audio data* where we slide the filter over *time*
- we could also have *3D convolutions*
 - ↳ The filter slides in *all three direction*, i.e., *left-to-right*, *up-to-down* and *front-to-back*
 - ↳ This is useful with *3D images*, e.g., 3D medical image of brain

Other Forms of Convolution



1D Convolution

2D Convolution



3D Convolution

There are also *other forms*, e.g., depth-wise convolution

at the end of day, they all *slide* over *input* with *some filter*